



Module 11

Synchronization



Overview

- Objectives
- Relevance



Process Synchronization

Signals	kill (2)	Sends a signal.
	sigaction (2)	Catches a signal.
Record locking	fcntl (2)	Locks part of a file. Exclusive read and write locks provided can be blocking or non-blocking.
System V semaphores	semget (2)	Gets a semaphore ID.
	semctl (2)	Sets a semaphore explicitly.
	semop (2)	Performs a set of semaphore operations (raise, lower, and wait).



Process Synchronization

Mechanism	Comment
mutex lock	For mutual exclusion; simple and fast.
reader-writer lock	Multiple-reader single-writer lock; useful when writes are fast and infrequent compared to reads.
semaphore	Pure Dijkstra semaphore.
condition variable	Useful in producer-consumer situations.

Thread Synchronization

- Coordination
- Single-thread processes
- Multithreaded processes
- Locking mechanisms
- Test and Set

Types of Synchronization Mechanisms

- Mutual exclusion lock
- Multiple-reader single-writer lock
- Semaphore
- Condition variable

Thread Synchronization Within a Process

- Multithreaded POSIX mutex locks
- Rwlocks and shared data structures
- Semaphores regulating bank teller resources
- Condition variables and a threadsafe stack

mutex.c

```
1  /*****
2  This program demonstrates the use of a POSIX thread
3  mutex to coordinate access to a shared resource.
4  Two threads are created, both of which access a
5  common counter. Delays are used to cause contention to
6  prove that the mutexes really work.
7
8  This program operates in two modes. With no
9  commandline arguments, the program operates without
10 synchronization and the threads fall all over
11 themselves. With argument "sync" passed on the
12 commandline, synchronization is used and the
13 program operates correctly.
14 *****/
15
16 #include <sys/types.h> /* For general. */
17 #include <stdio.h> /* Standard IO Library. */
18 #include <pthread.h> /* Posix threads. */
19 #include <errno.h> /* Error definitions. */
```



```

20
21 #include "utils.h"/* for fractSleep() and printWithTime(). */
22
23 #define LOOP_MAX 20/* Each thread counts up this many
iterations.*/
24
25
26
27 /* Global stuff. */
28
29 /* Common counter incremented by all threads. */
30 int commonCounter= 0;
31
32 /* mutex lock. Note the initializer can be used only
33 in this form of declaration. */
34 pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;
35
36 /* Delay to cause contention. */
37 double delay = 0.2;
38
39 /* Synchronized flag. */
40 boolean_t synchronized = B_FALSE;

```



```

41
42 /* Forward references. */
43 void * threadMain(void *);
44
45 /*****
46 main(int argc, char *argv[]) {
47 /*****
48 Main program. Get commandline arguments and spawn threads.
49 /*****
50
51 /* Define two threads. */
52 pthread_t threadID_1;
53 pthread_t threadID_2;
54
55 /* We have an argument and it is "sync": run in
synchronized mode.*/
56 if ((argc > 1) && (!(strcmp(argv[1], "sync")))) {
57     synchronized = B_TRUE;
58 }
59
60 printf ("\n%s running with sync %s\n\n", argv[0],
(synchronized) ? "Enabled" : "Disabled");
61

```



```

62
63 /* Create threads. */
64 if ((pthread_create( &threadID_1, NULL,
65     threadMain, NULL)) ||
66     (pthread_create( &threadID_2, NULL,
67     threadMain, NULL))) {
68     perror ("Error starting thread(s)");
69     exit (errno);
70 }
71
72 /* Wait for threads to finish. */
73 pthread_join(threadID_1, (void **)NULL);
74 pthread_join(threadID_2, (void **)NULL);
75 }
76
77 /*****
78 void * threadMain(void *dummy) {
79 /*****
80 This is the main for each thread. Get counter
81 value from global, hold it for a while to cause
82 contention, then replace it with an incremented
83 value. While it is held, another thread tries to

```



```

84 get it and increment it as well.
85
86 When running in synchronized mode, get the mutex
87 lock before accessing the global, and release it
88 only after we are done with the global. This causes
89 other threads to wait for the global until we are
90 finished with it.
91 /*****
92
93 int loopCount;/* Loop counter. */
94 int temp;/* Used in modification delay. */
95 char buffer[80];/* For writing status. */
96
97 for (loopCount = 0; loopCount < LOOP_MAX;
98     loopCount++){
99
100 /* Get mutex lock if running in synchronized
101 mode. */
102 if (synchronized) {
103     pthread_mutex_lock(&mutex);
104 }
105

```



```

106      /* Now that we have the lock, we can safely
107         increment the counter without worry that the
108         other thread accesses it at the same time. */
109
110      /* Print out initial value. */
111      sprintf (buffer, "Common counter incremented "
112              "from %d to ", commonCounter);
113      write(STDOUT_FILENO, buffer, strlen(buffer));
114
115      /* Modify common counter. Take a long time to
116         do it to cause contention. */
117      temp = commonCounter;
118      fractSleep(delay);
119      commonCounter = temp + 1;
120
121      /* Update status output. */
122      sprintf (buffer, "%d\n", commonCounter);
123      write(STDOUT_FILENO, buffer, strlen(buffer));
124
125      /* Release mutex lock if running in
126         synchronized mode. */
127      if (synchronized) {

```



```

128      pthread_mutex_unlock(&mutex);
129  }
130
131      /* Delay for half as long as the increment
132         delay, to cause contention. */
133      fractSleep(delay/2.0);
134  }
135  }
136  /*****/
137  void fractSleep(float time) {
138  /*****/
139      Routine to delay for fractions of a second.
140  *****/
141
142      static struct timespec timeSpec;
143
144      /* Truncate fraction in int conversion. */
145      timeSpec.tv_sec = (time_t)time;
146
147      /* Load fraction. */
148      timeSpec.tv_nsec = (int)((time - timeSpec.tv_sec)*
149                               1000000000);

```



```

149
150  /* Delay. */
151  nanosleep(&timeSpec, NULL);
152  }

```



rW.C

```

.
.
32
33  /* Multiple Reader Single Writer lock. */
34  pthread_rwlock_t rwlock =
35      PTHREAD_RWLOCK_INITIALIZER;
.
.
124  /*****/
125  This the main for the writer thread. Get the
126  appropriate lock, do the read or write, and then
127  release the lock. Note that trying to get either
128  lock can force the thread to block if the lock is
129  unavailable.
130  *****/
131  void * writerMain(void *arg) {
132      int loopCount;
133
134      for (loopCount = 0; loopCount < LOOP_MAX;
135           loopCount++) {

```



```
136
137     /* Get either the rwlock or the mutex,
138     depending on the mode. */
139     if (userwlock) {
140         pthread_rwlock_wrlock(&rwlock);
141     }else {
142         pthread_mutex_lock(&mutex);
143     }
144     doWrite();
145
146     /* Release the lock. */
147     if (userwlock) {
148         pthread_rwlock_unlock(&rwlock);
149     }else {
150         pthread_mutex_unlock(&mutex);
151     }
152 }
153 }
154 }
155
156 /*****
157 This is the main for the reader threads. Get the
```



```
158 locks as for the writer threads, except specify
159 that this is a reader and not a writer when getting
160 the multiple reader single writer lock.
161 /*****/
162 void * readerMain(void * idArg) {
163     int id = (int)idArg; /* Arg is thread ID number. */
164     int loopCount;
165
166     for (loopCount = 0; loopCount < LOOP_MAX;
167         loopCount++) {
168
169         /* Get either the rwlock or the mutex,
170         depending on the mode. */
171         if (userwlock) {
172             pthread_rwlock_rdlock(&rwlock);
173         }else {
174             pthread_mutex_lock(&mutex);
175         }
176
177         doRead();
178
179         /* Release the lock. */
```




```
180     if (userwlock) {
181         pthread_rwlock_unlock(&rwlock);
182     }else {
183         pthread_mutex_unlock(&mutex);
184     }
185 }
186 }
```



Thread Synchronization Within a Process

Condition Variables





stack.c

```

1  /*****
2  This module exemplifies a threadsafe stack. It uses
3  a mutex to coordinate access to shared data. It
4  also demonstrates the use of condition variables by
5  which threads wait for the shared data to change,
6  and by which threads notify/resume each other when
7  the state has changed.
8  *****/
9
10 #include <pthread.h> /* Posix threads. */
11
12 #include "stack.h" /* Stack function definitions. */
13 #include "utils.h" /* for printWithTime(). */
14
15 #define STACK_SIZE 3 /* Define a small stack size to
16                        cause contention. */
17
18 /* Define the data structure shared between the
19    threads. */
20

```



```

21 static char buffer[STACK_SIZE]; /* Stack's buffer */
22 static int index = 0; /* Stack's index. */
23
24 /* Mutex lock for exclusive data access. */
25 static pthread_mutex_t mutex =
26     PTHREAD_MUTEX_INITIALIZER;
27
28 /* Condition var for coordinating threads. */
29 static pthread_cond_t conditionVar =
30     PTHREAD_COND_INITIALIZER;
31
32 /*****
33 void push(char oneChar) {
34 /*****
35 Push a character onto the stack. Return in the
36 second argument the stack index corresponding to
37 where the character is pushed.
38 /*****
39 char string[25];
40
41 /* Get the lock before accessing the shared data

```



```

41 in any way. */
42 pthread_mutex_lock(&mutex);
43
44 /* Test the data while under mutex protection, to
45 see if the stack is pushable. Don't change the
46 data yet as the mutex lock is given up if
47 pthread_cond_wait() is called; a free lock
48 implies consistent data. */
49
50 while (index == STACK_SIZE) {
51     printWithTime ("push sleeping...\n");
52     pthread_cond_wait(&conditionVar, &mutex);
53 }
54
55 /* Stack is pushable. Push the data. */
56 buffer[index++] = oneChar;
57
58 sprintf (string, "Pushed:\tchar %c\tindex %d\n",
59         oneChar, index-1);
60 printWithTime (string);
61
62 /* Notify waiting threads (poppers in this case)

```



```

63 that the stack has changed and, due to the
64 operation just completed, there is now
65 something to pop. */
66 pthread_cond_signal(&conditionVar);
67
68 /* Release the mutex. All done with shared data
69 access. */
70 pthread_mutex_unlock(&mutex);
71 }
72
73 /*****
74 char pop() {
75 /*****
76 Pop a character from the stack. Return in the
77 second argument the stack index corresponding to
78 where the character is popped.
79 /*****
80 char toReturn;
81 char string[25];
82
83 /* Get the lock before accessing the shared data
84 in any way. */

```



```

85  pthread_mutex_lock(&mutex);
86
87  /* Test the data while under mutex protection
88     to see if the stack is pushable. Don't change
89     the data yet as the mutex lock is given up if
90     pthread_cond_wait() is called; a free lock
91     implies consistent data. */
92
93  while (index == 0) {
94      printWithTime ("pop sleeping...\n");
95      pthread_cond_wait(&conditionVar, &mutex);
96  }
97
98  /* Stack is poppable. Pop the data. */
99  toReturn = buffer[--index];
100
101  sprintf (string, "Pop:\tchar %c\tindex %d\n",
102          toReturn, index);
103  printWithTime(string);
104
105  /* Notify waiting threads (pushers in this case)
106     that the stack has changed and, due to the

```



```

107      operation just completed, there is now something
108      to push. */
109  pthread_cond_signal(&conditionVar);
110
111  /* Release the mutex. All done with shared data
112     access. */
113  pthread_mutex_unlock(&mutex);
114
115  return toReturn;
116 }

```



proc_sem.c

```

.
.
53 /* This is the semaphore which regulates the bank
54    tellers. It operates like a mutex but can manage
55    more than one resource. */
56 sem_t bankLine;
57
58 pthread_mutex_t semMutex =
59     THREAD_MUTEX_INITIALIZER;
60 pthread_cond_t semCond = PTHREAD_COND_INITIALIZER;
61
62

```



```

80 /* Initialize semaphore to manage NUM_TELLERS
81    quantity of a resource. Set up so that only
82    threads of this process can use it. */
83 if (sem_init(&bankLine, THREADS_OF_THIS_PROCESS,
84             NUM_TELLERS)) {
85     perror ("Error initializing semaphore");
86     exit (errno);
87 }
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126 /* This is one of those rushed clients. */
127
128
129
130 if (inaHurry) {
131
132     /* See if a teller is available. If not,
133        sem_trywait returns EAGAIN. */
134     if ((sem_trywait(&bankLine) == -1)
135         && (errno == EAGAIN)) {
136
137         ...
138     }
139 }

```



```

145  /* If get here, there is a teller available right
146     away. Drop thru and do business. */
147  } else {
148
149  #if REPORT_NUM_SEM_FREE
150      sem_getvalue(&bankLine, &availableTellers);
151      pthread_mutex_unlock(&semMutex);
152      sprintf (string,
153              "Client %d is in a hurry and got a
154              teller. Count is now %d\n",
155              whoAmI, availableTellers);
156      printWithTime(string);
157  #else
158      sprintf (string,
159              "Client %d%c is in a hurry and got a
160              teller.\n", whoAmI);
161      printWithTime(string);
162  #endif
161  }
162

```



```

163  /* Not in a hurry. */
164  } else {
165      /* Wait for next available teller. */
166
167  #if REPORT_NUM_SEM_FREE
168
169      /* Calling sem_trywait() in a loop until we get
170         the lock, and sleeping on the condition
171         variable if the lock is unavailable is
172         necessary to prevent deadlock, since we have
173         the mutex. */
174
175      while ((sem_trywait(&bankLine) == -1)
176             && (errno == EAGAIN)) {
177          pthread_cond_wait(&semCond, &semMutex);
178      }
179      sem_getvalue(&bankLine, &availableTellers);
180      sprintf (string, "Client %d got a teller.
181                      Count is now %d\n", whoAmI,
182                      availableTellers);
183      printWithTime(string);
184  #endif
185  }
186

```



```

180      pthread_mutex_unlock(&semMutex);
181
182  #else
183
184      sem_wait(&bankLine);
185      sprintf (string, "Client %d got a teller.\n",
186              whoAmI);
187      printWithTime(string);
188  #endif
189  }
190
191  /* We've got a teller. Do our business. */
192  doBusiness(whoAmI);
193
194  #if REPORT_NUM_SEM_FREE
195      pthread_mutex_lock(&semMutex);
196
197      /* Done doing business. Teller is available. */
198      sem_post(&bankLine);
199

```



```

200  /* Get value of the semaphore to report back as a
201     diagnostic. */
202  sem_getvalue(&bankLine, &availableTellers);
203
204  pthread_cond_signal(&semCond);
205  pthread_mutex_unlock(&semMutex);
206
207  sprintf (string,
208          "Client %d done doing business. Available"
209          "teller count: %d\n", whoAmI, availableTellers);
210  printWithTime(string);
211  #else
212
213      /* Done doing business. Teller is available. */
214      sem_post(&bankLine);
215

```




Using Interprocess Synchronization

- Shared area setup
- Mutex lock initialization
- Cleanup



Using Interprocess Synchronization

```
.
.
14  /* Shared data is created in this mmaped file. */
15  #define SHARED_DATA_FILE "/tmp/proc_mutex.dat"
16
17  /* Global stuff. */
18
19  typedef struct globalStuff {
20      pthread_mutex_t sharedMutex;
21      int count;
22  } globalStuff;
23
24  /* Handle to shared area. */
25  struct globalStuff * globalArea;
26
27
```



```
28  /*****
29  void cleanup(int x) {
30  /*****
31  Cleanup handler, called as a signal handler for ^C.
32  Unlock the mutex before terminating.
33  /*****
34      printf ("Cleaning up...\n");
35      pthread_mutex_unlock(&globalArea->sharedMutex);
36      exit(0);
37  }
.
.
48  /* For interprocess mutex init. */
49  pthread_mutexattr_t mutex_attributes;
50  /* True if we set up shared area. */
51  boolean_t createdHere = B_FALSE;
.
.
```



```
52
53  /* Open the shared area. Err out if it already
54      exists (which means we are
55      not the first process to use it, and thus not
56      to be the creator if it. */
57  shared_fd = open(SHARED_DATA_FILE, (O_RDWR |
58      O_CREAT | O_EXCL), 0777);
59
60  /* Flag that we are the creator if we are, or
61      just open the file otherwise. */
62  if ((shared_fd == -1) && (errno == EEXIST)) {
63      shared_fd = open(SHARED_DATA_FILE,
64      O_RDWR, 0777);
65  } else {
66      createdHere = B_TRUE;
67  }
68
69  /* Some other error opening the file. */
70  if (shared_fd < 0) {
71      perror ("Error opening shared mutex area");
72      exit (errno);
73  }
```



```
69
70  /* Expand the file to hold the shared data,
    before mmaping it. */
71  if (createdHere) {
72      if (ftruncate(shared_fd,
73          sizeof(struct globalStuff)) {
74          perror ("ftruncate");
75          exit (errno);
76      }
77  }
78  /* Mmap the file so all processes needing access
    to it can see it. */
79  globalArea = (struct globalStuff *)mmap(
80      NULL, sizeof(struct globalStuff), PROT_READ
81      | PROT_WRITE, MAP_SHARED,
82      shared_fd, 0);
83  if (globalArea == (struct globalStuff *)-1) {
84      perror ("Error getting shared Mutex virtual
85      addr");
86      exit (errno);
87  }
```



```
86
87  /* Set up cleanup handler for ^C. */
88  signal(SIGINT, cleanup);
89
90  /* Initialize mutex in shared region if we are
    the shared region creator. */
91  if (createdHere) {
92
93      /* Initialize mutex attributes list. */
94      if (pthread_mutexattr_init(&mutex_attributes)) {
95          perror ("pthread_mutexattr_init");
96          exit (errno);
97      }
98
99      /* Reflect in the attributes that this is an
    interprocess shared mutex. */
100     if (pthread_mutexattr_setpshared(
101         &mutex_attributes,
102         PTHREAD_PROCESS_SHARED)) {
103         perror ("pthread_mutexattr_setpshared");
104         exit (errno);
105     }
```



```
105
106  /* Initialize the mutex in the mmaped area,
    with the attributes above. */
107  if (pthread_mutex_init(&globalArea-
108      >sharedMutex, &mutex_attributes)) {
109      perror ("pthread_mutex_init");
110      exit (errno);
111  }
112  .
113  .
```



System V IPC Semaphores

Command	Functionality	Argument
IPC_RMID	Remove group	Ignored
SETVAL	Set value of particular semaphore	int value*
GETVAL	Get value of particular semaphore	Ignored**
SETALL	Set value of each semaphore in group	ushort_t * array of values*
GETALL	Get value of each semaphore in group	ushort_t * array of values*

* Interpretation of union semun argument.
** Ignored as value is returned in function return.



ipcsem.c

```

1  /*****
2  This program demonstrates the use of an IPC semaphore
3  in coordinating tasks which require simultaneous
4  multiple resource allocations, each requesting more
   than one of the resource.
5
6  Allocate a semaphore to represent the resources of a
7  moving company with five trucks, 12 human movers,
8  and an insurance policy of $1 Million to be spread
   across all concurrent jobs.
9
10 Jobs are done on a first come, first served basis.
11 A job goes only if there are an adequate number of
12 movers, trucks and insurance to cover the job,
13 after all other jobs are accounted for. If the job
14 is a go, decrement the resource count, so that all
   resources in use are accounted for when the next
   job is requested.
15 *****/
16

```



```

17 #include <sys/types.h> /* For general. */
18 #include <sys/ipc.h> /* System 5 IPC defs. */
19 #include <sys/sem.h> /* System 5 IPC semaphore
   defs. */
20 #include <pthread.h> /* Posix threads. */
21 #include <stdlib.h> /* Needed for delay to work
   properly. */
22
23 #include "utils.h" /* For fractSleep() and
   printWithTime(). */
24
25 #define NUM_THREADS 10 /* Number of simultaneous
   requests. */
26 #define TIME_BTWN_NEW_THREADS 0.5 /* Time between
   intro of new request. */
27 #define RUNTIME_RANGE 5.0 /* Time of longest job. */
28

```



```

29 /* These define a single semaphore group consisting
30 of three semaphores. All the semaphores in a
31 group can be modified together in a single atomic
32 operation. The first semaphore represents the
33 number of trucks available; the second: the
   number of movers available; the third: the amount
   of insurance in $1000 units available. */
34
35 #define NUM_SEMS_IN_GROUP 3
36 #define TRUCK_SEM 0
37 #define MOVER_SEM 1
38 #define INSUR_SEM 2
39
40 /* These are used to initialize semaphores with
   total resources managed. */
41 #define NUM_TRUCKS 5
42 #define NUM_MOVERS 12
43 #define AMT_INSUR 1000
44

```



```

45 /* Flags passed thru sembuf structure; wait for
   resource. */
46 #define WAIT 0
47
48 #define STRING_SIZE 80
49
50 #define FALSE 0
51 #define TRUE (!FALSE)
52
53
54
55 /* Define some jobs here. This table specifies how
56 many of each resource is requested per each job.
57 The values in this table cause contention: for
58 example, the first job uses 4 trucks out of 5
59 available, making some of the others wait for a
   truck. */
60
61

```



```

62 struct job { int numTrucks; int numMovers;
               int amtInsurance; } jobTable[] = {
63     { 4,5,250},/* Use many trucks */
64     { 1,2,500},/* Normal amts. */
65     { 3,5,1000},/* Lots of insur */
66     { 2,8,250},/* Many movers. */
67 };
68
69 /* Number of jobs in the table. */
70 int numJobs = sizeof(jobTable) /
               sizeof (struct job);
71
72 int semid;/* IPC semaphore identifier (semaphore
               accessed through this).*/
73
74 extern int errno;
75
76 void *threadMain(void *);
77
78

```



```

79 /*****
80 main() {
81 /*****
82 Main function. Allocate and initialize resources,
83 then spawn threads.
84 /*****
85
86 /* Array of threads, one per request. */
87 pthread_t threads[NUM_THREADS];
88
89 /* Needed to convert constant for semctl call.*/
90 int numTrucks = NUM_TRUCKS;
91 int numMovers = NUM_MOVERS;
92 int amtInsur= AMT_INSUR;
93
94 int count;
95
96 /* Allocate a single semaphore group with three
97    semaphores in it. */

```



```

98 if ((semid = semget(IPC_PRIVATE,
99     NUM_SEMS_IN_GROUP, IPC_CREAT | 0600)) == -1){
100     perror ("semget");
101     exit (errno);
102 }
103
104 /* Initialize each semaphore in the group.
105    Could also have used SETALL. */
106 if ((semctl(semid, TRUCK_SEM, SETVAL,
107     &numTrucks)) || (semctl(semid,
108     MOVER_SEM, SETVAL, &numMovers)) ||
109     (semctl(semid, INSUR_SEM, SETVAL,
110     &amtInsur))) {
111     perror ("Error initializing semaphores");
112     goto cleanup;
113 }
114
115 /* Initialize random number generator used for time
116    delays. */
117 srand48 (time(NULL));
118

```



```

116 /* Spawn the threads. The argument passed to
117    threadMain is a job table index, so multiple
118    requests can be made using the same table entry
119    when the index wraps. Delay to model jobs being
120    staggered instead of coming all at once. */
121 for (count = 0; count < NUM_THREADS; count++) {
122     if (pthread_create(
123         &threads[count], NULL, threadMain,
124         (void *) (count % numJobs))) {
125         perror ("Error starting reader threads");
126         goto cleanup;
127     }
128     fractSleep(TIME_BTWN_NEW_THREADS);
129 }
130
131 /* Wait for threads to finish. */
132 for (count = 0; count < NUM_THREADS; count++) {
133     pthread_join(threads[count], (void **)NULL);
134 }
135

```



```

136 cleanup:
137
138 /* Delete the semaphore. This is not done
139    automatically by the system. */
140 if (semctl(semid, 0, IPC_RMID, NULL)) {
141     perror ("semctl IPC_RMID:");
142 }
143 }
144
145
146 /*****
147 void *threadMain(void * arg) {
148     *****/
149 Here is where a thread starts executing. A
150 thread in executing this function represents a
151 task for the moving company, and requests the
152 resources it needs.
153 /*****/
154

```



```

155 /* The argument passed in is a table index. */
156 int jobNum = (int)arg;
157
158 /* Local string for message composition. */
159 char string[STRING_SIZE];
160
161 sprintf (string,
162         "Job # %d requesting %d trucks, %d people, "
163         "$%d000 insurance...\n",
164         jobNum, jobTable[jobNum].numTrucks,
165         jobTable[jobNum].numMovers,
166         jobTable[jobNum].amtInsurance);
167 printWithTime(string);
168
169 /* Get the resources needed. Wait for them if
170    necessary. */
171 if (reserve(semid, jobTable[jobNum])) {
172     perror ("reserve");
173     return (NULL);
174 }
175

```



```

176 sprintf (string,
177         "Job # %d got %d trucks, %d people"
178         ", %d000 insurance and is running...\n",
179         jobNum, jobTable[jobNum].numTrucks,
180         jobTable[jobNum].numMovers,
181         jobTable[jobNum].amtInsurance);
182 printWithTime(string);
183
184 /* Delay to simulate the time the resources are
185    in use. */
186 fractSleep(drand48() * RUNTIME_RANGE);
187
188 sprintf (string,
189         "Job # %d done; returning %d "
190         "trucks, %d people, %d000 insurance...\n",
191         jobNum, jobTable[jobNum].numTrucks,
192         jobTable[jobNum].numMovers,
193         jobTable[jobNum].amtInsurance);
194 printWithTime(string);
195

```



```

194 /* Release resources. */
195 if (release(semid, jobTable[jobNum])) {
196     perror ("release");
197 }
198 }
199
200
201 /*****
202 int reserve(int semid, struct job thisJob) {
203     *****/
204 Reserve resources required for a job. This
205 wrapper function passes negative values into
206 playWithSemaphores(), since negative values
207 represent resource allocations.
208 /*****/
209 return (playWithSemaphores(
210         semid, -thisJob.numTrucks, -thisJob.numMovers,
211         -thisJob.amtInsurance));
212 }
213

```



```
214 /*****  
215 int release(int semid, struct job thisJob) {  
216 /*****  
217     Release resources of a job completed. This  
218     wrapper function passes positive values into  
219     playWithSemaphores(), since positive values  
220     represent resource deallocations.  
221 /*****  
222     return (playWithSemaphores(  
223         semid, thisJob.numTrucks, thisJob.numMovers,  
224         thisJob.amtInsurance));  
225 }  
226  
227  
228 /*****  
229 static int playWithSemaphores(  
230     int semid, int numTrucks, int numMovers,  
231     int amtInsurance) {  
232 /*****  
233     This is the workhorse function that allocates /  
234     deallocates resources from the semaphore group.  
235 /*****
```



```
236  
237 /* There is one operation per semaphore for this  
238     example. This allocates an array of semaphore  
239     operations, all of which is carried out with a  
240     single atomic operation (system call to  
241     semop()).*/  
242 struct sembuf ops[NUM_SEMS_IN_GROUP];  
243  
244 /* One operation per semaphore. Note that a  
245     negative value to ops[x].sem_op allocates a  
246     resource, while a positive value releases it. */  
247  
248 ops[0].sem_num = TRUCK_SEM;  
249 ops[0].sem_op = numTrucks;  
250 ops[1].sem_num = MOVER_SEM;  
251 ops[1].sem_op = numMovers;  
252 ops[2].sem_num = INSUR_SEM;  
253 ops[2].sem_op = amtInsurance;  
254
```




```
255 /* All semaphore operations are to be handled in  
256     the same way. */  
257 ops[0].sem_flg = ops[1].sem_flg =  
258     ops[2].sem_flg = WAIT;  
259  
260 /* "The call" that does the work. */  
261 return (semop(semid, ops, NUM_SEMS_IN_GROUP));  
262 }  
263
```



Exercise: Synchronization

- Objectives
- Tasks
- Discussion
- Solutions



Module 12

Sockets