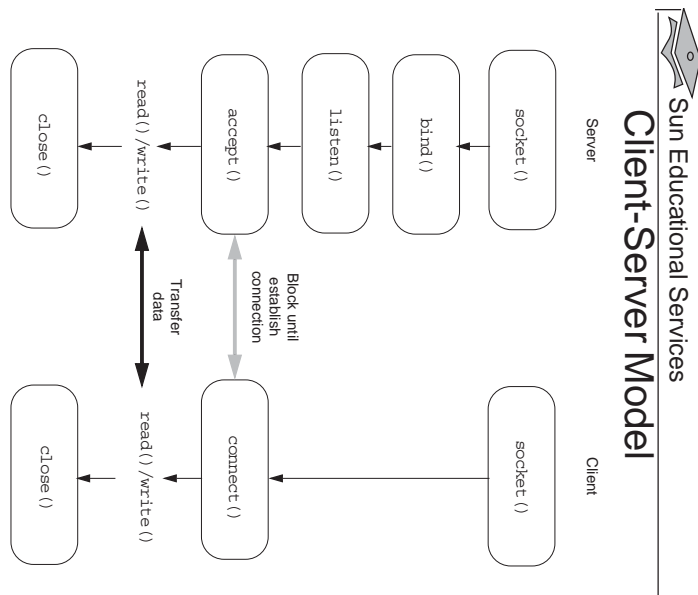


# Module 12

## Sockets

## Overview

- Objectives
- Relevance



## Socket Types

### SOCK\_STREAM – TCP

- Connection-oriented
- No message boundaries
- Reliable
- Sequenced
- Easier to use with reliability, but more expensive

### SOCK\_DGRAM – UDP

- Connectionless
- Message boundaries
- Not reliable, not sequenced
- Duplicates possible
- More efficient

## Combining Socket Domains and Types

- AF\_UNIX, SOCK\_STREAM
- AF\_UNIX, SOCK\_DGRAM
- AF\_INET, SOCK\_STREAM
- AF\_INET, SOCK\_DGRAM

## Combining Socket Domains and Types

### Owner's Activities

Human Speak	Computer Speak	System Call
Buy a phone	Establish an end point	socket (3SOCKET)
Get a phone number	Establish a rendezvous	bind (3SOCKET)
Activate the phone line.	Set queue length and enable service	listen (3SOCKET)
Wait for a new client and redirect to Cesare's phone.	Acknowledge	accept (3SOCKET) and fork (2)
Cesare takes order	Full duplex conversation	read (2), write (2), recv (3SOCKET), and send (3SOCKET)

## Combining Socket Domains and Types

### Client's Activities

Human Speak	Computer Speak	System Call
Walk to any phone.	Establish an end point	socket ()
Look for "Pizza" in phone book	Use directory services	gethostbyname (3NSL)
Dial the number	Request connection	connect (3SOCKET)
Place order	Full duplex conversation	read (), write (), recv (), and send ()
Hang up	Send EOF	close (2)

## Combining Socket Domains and Types

### AF\_UNIX, SOCK\_STREAM Call Sequences

Server	Client
sd = socket	sd = socket
bind	
listen	
nsd = accept	connect
write, read, send, recv	write, read, send, recv
close (nsd)	close (sd)
close (sd)	



## Combining Socket Domains and Types

AF\_UNIX, SOCK\_DGRAM Call Sequences

Server	Client
sd = socket	sd = socket
bind	bind
recvfrom, sendto	sendto, recvfrom
close(sd)	close(sd)



## Combining Socket Domains and Types

AF\_INET, SOCK\_STREAM Call Sequences

Server	Client
sd = socket	sd = socket
bind	
listen	gethostbyname
nsd = accept	connect
write, read, send, recv	write, read, send, recv
close(nsd)	close(sd)
close(sd)	



## Combining Socket Domains and Types

AF\_INET, SOCK\_DGRAM Call Sequences

Server	Client
sd = socket	sd = socket
bind	bind
	gethostbyname
recvfrom, sendto	sendto, recvfrom
close(sd)	close(sd)



## Creating and Destroying Sockets

- socket () – Creates a socket
- shutdown () – Destroys a socket



## Binding an Address to a Socket

- `bind()` – Assigns a local address (or name) to the socket descriptor
- `listen()` – Activates a socket



## Binding an Address to a Socket

### UNIX Example

```
1  int  sd;
2  struct sockaddr_un name;
3
4  sd = socket(AF_UNIX, SOCK_STREAM, 0);
5  name.sun_family = AF_UNIX;
6  strcpy(name.sun_path, "/tmp/socket1");
7  unlink( name.sun_path ); /* Otherwise bind could fail */
8
9  if( bind(sd, (struct sockaddr*)&name, sizeof(name) ) == -1 )
10 {
11     perror("bind");
12     exit(1);
13 }
```



## Binding an Address to a Socket

### Internet Example

```
1  struct sockaddr_in sin;
2  int  sd;
3
4  sd = socket(AF_INET, SOCK_STREAM, 0);
5  memset((char *)&sin, '\0', sizeof(sin));
6  sin.sin_family = AF_INET;
7  sin.sin_port = htons(7000); /* pick */
8  sin.sin_addr.s_addr = htonl(INADDR_ANY);
9  if ( bind(sd, (struct sockaddr *)&sin, sizeof(sin))
10      == -1 ) {
11      perror("bind");
12      exit(1);
13  }
14
15  if ( listen(sd, 5) == -1 ) {
16      perror("listen?");
17      exit(1);
18  }
```



## Accepting a Connection

```
1  struct sockaddr_un client;
2  int  len = sizeof(client), comm;
3  ....
4
5  loop:
6  while ((comm = accept(sd, &client, &len)) != -1) {
7      if (fork() == 0) {
8          /* child to handle session */
9          close(sd); /* do not need rendezvous */
10         do_service(comm); /* talk to client using
11                             comm */
12         close(comm);
13         exit(0);
14     }
15     close(comm); /* do not need communication */
16 }
```



```
16  /* If a signal arrives during accept(), it fails
17  with -1. So, if catching signals, use the following
18  2 lines */
19  if (errno == EINTR) {
20      goto loop;
21  }
```



## Connecting to the Server

### Client Side Example

```
1  int sd;
2  struct sockaddr_un name;
3
4  sd = socket(AF_UNIX, SOCK_STREAM, 0);
5  name.sun_family = AF_UNIX;
6  strcpy( name.sun_path, "/tmp/socket1");
7
8  if (connect(sd, (struct sockaddr *)&name,
9             sizeof(name)) == -1) {
10     perror("connect");
11     exit(1);
12 }
```



## Connecting to the Server

### Internet Socket Example

```
1  int sd;
2  struct sockaddr_in sin;
3  struct hostent *he;
4
5  sd = socket(AF_INET, SOCK_STREAM, 0);
6
7  memset((char *)&sin, '\0', sizeof(sin));
8  sin.sin_family = AF_INET;
9
10 /* match server */
11 he = gethostbyname("Server");
12 sin.sin_port = htons(7000);
13 memcpy((char *)&sin.sin_addr, he->h_addr_list[0],
14        he->h_length);
15
```



```
16 if (connect(sd, (struct sockaddr *)&sin,
17             sizeof(sin)) == -1) {
18     perror("connect");
19     exit(1);
20 }
```

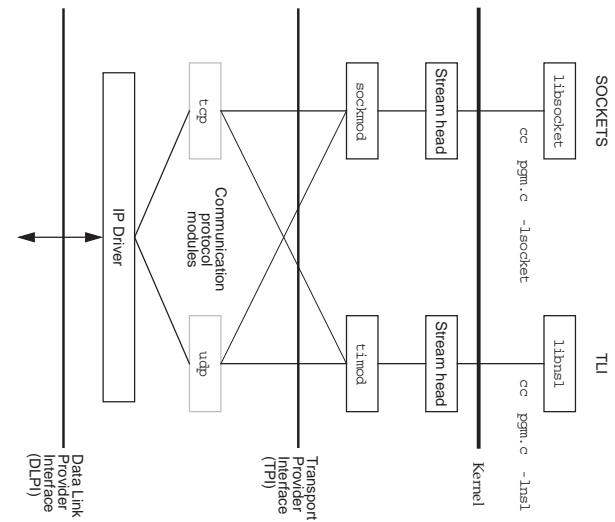


## Receiving Data

```
1 char message[80];
2 int length = sizeof(message);
3 int rval, i;
4
5 if ((rval = recv(sd, message, length, 0)) == -1) {
6     perror("recv");
7     exit(1);
8 }
9 for (i = 0; i < rval; i++) {
10     putchar(message[i]);
11 }
```



## Network Sockets and TLI



## Network Sockets and TLI

	Sockets	TLI
Performance	Comparable – both use STREAMS	Comparable – both use STREAMS
POSIX	Accepted for adoption	Accepted for adoption
Transport independence	No	Yes, but must purchase unbundled software if you want to support other than TCP and UDP



## Establishing a Connect to a Remote Host

```
1 /* gethost.c */
2
3 #include <sys/types.h>
4 #include <sys/socket.h>
5 #include <netdb.h>
6 #include <stdio.h>
7
8 main(int argc, char *argv[]){
9
10     struct hostent *he;
11     char **tmp;
12
13     if (argc != 2) {
14         fprintf(stderr, "Usage: %s hostname\n", argv[0]);
15         exit(1);
16     }
17     he = gethostbyname(argv[1]);
18
19     if (he == NULL) {
```



```

20     fprintf(stderr, "No such host.\n");
21     exit(1);
22 }
23 printf("Official name of %s is %s.\n",
24        argv[1], he->h_name);
25
26 for (tmp=he->h_aliases; *tmp != NULL; tmp++) {
27     printf("\tAlias: %s\n", *tmp);
28 }
29
30 tmp = he->h_addr_list;
31 printf("IP address: %d.%d.%d.%d\n",
32        (unsigned char) (*tmp)[0],
33        (unsigned char) (*tmp)[1],
34        (unsigned char) (*tmp)[2],
35        (unsigned char) (*tmp)[3]);
36 return 0;
37 }

```



## Determining File Descriptor Activity

```

1  #include <sys/types.h>
2  #include <sys/socket.h>
3  #include <netinet/in.h>
4  #include <poll.h>
5  #include <errno.h>
6
7  /* Number of file descriptors to poll. */
8  #define NUM_FDS_TO_POLL 2
9
10 main() {
11
12     int udpsd1, udpsd2;
13     struct sockaddr_in sin1;
14     struct sockaddr_in sin2;
15     struct sockaddr_in returnAddr;
16     int returnAddrSize;
17

```



```

18  /* This array of struct pollfd defines fully
19  what poll() is to wait on, and provides a place
20  for poll() to respond. */
21  struct pollfd fds[NUM_FDS_TO_POLL];
22
23  /* Set up two UDP sockets to wait on. */
24
25  memset((char *)&sin1, '\0', sizeof(sin1));
26  sin1.sin_family = AF_INET;
27  sin1.sin_port = htons(7000);
28  sin1.sin_addr.s_addr = htonl(INADDR_ANY);
29
30  /* create one udp sd */
31  udpsd1 = socket(AF_INET, SOCK_DGRAM, 0);
32  if (udpsd1 == -1) {
33      perror("udp socket");
34      exit (errno);
35  }
36

```



```

37  if (bind(udpsd1, (struct sockaddr *)&sin1,
38          sizeof(sin1)) == -1) {
39      perror ("tcp bind");
40      exit (errno);
41  }
42
43  memset((char *)&sin2, '\0', sizeof(sin2));
44  sin2.sin_family = AF_INET;
45  sin2.sin_port = htons(7001);
46  sin2.sin_addr.s_addr = htonl(INADDR_ANY);
47
48  /* create one udp sd */
49  udpsd2 = socket(AF_INET, SOCK_DGRAM, 0);
50  if (udpsd2 == -1) {
51      perror ("udp socket");
52      exit (errno);
53  }
54

```



```
55  if (bind(udp2, (struct sockaddr *)&sin2,
56      sizeof(sin2)) == -1) {
57      perror ("udp bind");
58      exit (errno);
59  }
60
61  /*
62  Initialize each element of the fds array with:
63  - fd: the file descriptor of interest.
64  - events: what type of event to wait on.
65
66  Poll returns with the revents field set with the
67  events which are ready to do. E.G. if it
68  returns POLLIN, a read can be done without
69  blocking.
70  */
71
72  fds[0].fd = udp1;
73  fds[0].events = POLLIN;
74  fds[1].fd = udp2;
75  fds[1].events = POLLIN;
76
```



```
77  /*
78  Timeout value (poll's third argument) can be:
79  0: Immediate return whether or not fd is ready.
80  -1 (INFTIM): Blocks until a fd is ready.
81  <other value>: time in mS to wait for a fd
82  before returning.
83  */
84
85  for (;;) {
86      int numfds;
87      if ((numfds = poll(fds, NUM_FDS_TO_POLL, -1)) < 0) {
88          perror ("poll");
89          exit (errno);
90      } else {
91          printf ("Poll returned %d fds to read.\n",
92              numfds);
93      }
94
```



```
95  /* if udp sd ready, use it */
96  if (fds[0].revents == POLLIN) {
97      char buffer[80];
98      if (recvfrom(udp1, buffer, 80, 0,
99          (struct sockaddr *)&returnAddr,
100          &returnAddrSize) >= 0) {
101          puts(buffer);
102      } else {
103          perror ("recvfrom");
104      }
105  }
106
```



```
107  /* if udp sd ready, use it */
108  if (fds[1].revents == POLLIN) {
109      char buffer[80];
110      if (recvfrom(udp2, buffer, 80, 0,
111          (struct sockaddr *)&returnAddr,
112          &returnAddrSize) >= 0) {
113          puts(buffer);
114      } else {
115          perror ("recvfrom");
116      }
117  }
118  }
119  }
120
```





## Asynchronous I/O

```

1  int char_present;
2
3  void handler() {
4      char_present++;
5  }
6
7  int fd, flags;
8
9  /* Install handler for SIGPOLL, see Signals module*/
10     install_disp(SIGPOLL, handler);
11
12 /* Register fd to have SIGPOLL sent when IO is
13  * possible. Usually fd is a pipe or socket. */
14     ioctl(fd, I_SETSIG, S_RDNORM);
15
16     while(1) {
17

```



```

18 /* Block SIGPOLL to prevent char_present being changed */
19 if (char_present) {
20
21     /* Can use poll(2) to determine which
22     descriptor sent the SIGPOLL, if more than one. */
23     do_io();
24     char_present = 0;
25 }
26
27 /* Unblock SIGPOLL */
28 .
29 /* Main program work is in this loop */
30 .
31 }

```



## AF\_UNIX, SOCK\_STREAM Example

### Server Side

```

1  /* srvr1.c - AF_UNIX, SOCK_STREAM
2  * Server reads a value from client and
3  * sends back twice that value */
4
5  #include <stdio.h>
6  #include <sys/types.h>
7  #include <sys/socket.h>
8  #include <sys/un.h>
9
10     main() {
11
12         int sd, comm;
13         struct sockaddr_un myname, client;
14         int namesize = sizeof(struct sockaddr_un);
15         int clientsize = sizeof(struct sockaddr_un);
16         double dub;
17

```



```

18 /* Create server's rendezvous socket sd */
19 if ( (sd = socket(AF_UNIX, SOCK_STREAM, 0)) == -1) {
20     perror("srvr1.c:main:socket");
21     exit(1);
22 }
23
24 /* Fill in server's address and bind it to sd */
25 myname.sun_family = AF_UNIX;
26 strcpy(myname.sun_path, "/tmp/socket1");
27 unlink( myname.sun_path );
28
29 if( bind(sd, (struct sockaddr*)&myname, namesize)
30     == -1 ) {
31     perror("srvr1.c:main:bind");
32     exit(1);
33 }
34
35 /* Prepare to receive multiple connect requests */
36 if( listen(sd, 128) == -1 ) {
37     perror("srvr1.c:main:listen");
38     exit(1);
39 }

```



```

40
41  /* Infinite loop to accept client requests */
42  while(1) {
43      comm = accept(sd, (struct sockaddr*)&client,
44                  &clientsize);
45      if( comm == -1) {
46          perror("srvr1.c:main:accept");
47          exit(1);
48      }
49      read( comm, &dub, sizeof(dub));
50      dub += dub;
51      write( comm, &dub, sizeof(dub));
52      close(comm);
53  }
54
55  /* Install a signal handler for cleanup including:
56  * close(sd);
57  * unlink( myname.sun_path ); */
58  }

```



## Client Side

```

1  /* clnt1.c - AF_UNIX, SOCK_STREAM */
2
3  #include <stdio.h>
4  #include <sys/types.h>
5  #include <sys/socket.h>
6  #include <sys/un.h>
7
8  main() {
9
10     int namesize = sizeof(struct sockaddr_un);
11     int sd;
12     struct sockaddr_un srvr;
13     double dub=2.3, dub2;
14
15     /* Create client's socket sd */
16     if( (sd = socket(AF_UNIX, SOCK_STREAM, 0)) == -1) {
17         perror("clnt1.c:main:socket");
18         exit(1);
19     }
20

```



```

21  /* Fill in server's address and connect to server */
22  srvr.sun_family = AF_UNIX;
23  strcpy(srvr.sun_path, "/tmp/socket1");
24
25  if( connect(sd, (struct sockaddr*)&srvr, namesize)
26      == -1 ) {
27      perror("clnt1.c:main:connect");
28      exit(1);
29  }
30
31  /* Communicate with server */
32  write(sd, &dub, sizeof(dub));
33  read( sd, &dub2, sizeof(dub2));
34  close(sd);
35  printf("dub: %.2f, dub2: %.2f\n", dub, dub2);
36  }

```



## AF\_UNIX, SOCK\_DGRAM Example

## Server Side

```

1  /* srvr2.c - AF_UNIX, SOCK_DGRAM
2  * Server reads a string from client and
3  * sends back a greeting appended to the string */
4
5  #include <sys/types.h>
6  #include <sys/socket.h>
7  #include <sys/un.h>
8
9  #define BUFSZ  256
10
11  main() {
12
13      int sd, comm;
14      struct sockaddr_un myname, client;
15      int namesize = sizeof(struct sockaddr_un);
16      char buf[BUFSZ];
17

```



```

18  /* Create server's socket sd */
19  if( (sd = socket(AF_UNIX, SOCK_DGRAM, 0)) == -1) {
20      perror("srvr2.c:main:socket");
21      exit(1);
22  }
23
24  /* Fill in server's address and bind it to sd */
25  myname.sun_family = AF_UNIX;
26  strcpy(myname.sun_path, "/tmp/socket2srvr");
27  unlink( myname.sun_path );
28
29  if( bind(sd, (struct sockaddr*)&myname, namesize)
30      == -1 ) {
31      perror("srvr2.c:main:bind");
32      exit(1);
33  }
34

```



```

35  /* Infinite loop to accept client requests */
36  while(1){
37      if( recvfrom(sd, (char*)&buf, BUFSZ, 0,
38          (struct sockaddr*)&client, &namesize) == -1) {
39          perror("srvr2.c:main:recvfrom");
40          exit(1);
41      }
42      strcat(buf, ", Hello from the server" );
43
44      if( sendto(sd, buf, BUFSZ, 0,
45          (struct sockaddr*)&client, namesize) == -1) {
46          perror("srvr2.c:main:sendto");
47          exit(1);
48      }
49  }
50
51  /* Install a signal handler for cleanup including:
52  * close(sd);
53  * unlink( myname.sun_path ); */
54  }

```



## Client Side

```

1  /* clnt2.c - AF_UNIX, SOCK_DGRAM */
2
3  #include <sys/types.h>
4  #include <sys/socket.h>
5  #include <sys/un.h>
6
7  #define BUFSZ 256
8
9  main() {
10
11      int namesize = sizeof(struct sockaddr_un);
12      int sd;
13      struct sockaddr_un clnt, srvr;
14      char buf[BUFSZ];
15
16      /* Create client's socket sd */
17      if( (sd = socket(AF_UNIX, SOCK_DGRAM, 0)) == -1) {
18          perror("clnt2.c:main:socket");
19          exit(1);
20      }

```



```

21
22      /* Fill in client's address and bind it to sd */
23      clnt.sun_family = AF_UNIX;
24      strcpy(clnt.sun_path, "/tmp/socket2clnt");
25      unlink( clnt.sun_path );
26
27      if( bind(sd, (struct sockaddr*)&clnt, namesize)
28          == -1 ) {
29          perror("clnt2.c:main:bind");
30          exit(1); }
31
32      /* Fill in server's address for sendto server */
33      srvr.sun_family = AF_UNIX;
34      strcpy(srvr.sun_path, "/tmp/socket2srvr");
35
36      /* Communicate with server */
37      if( sendto(sd, "Deac", 5, 0, (struct sockaddr*)&srvr,
38          namesize) == -1) {
39          perror("clnt2.c:main:sendto");
40          exit(1);
41      }
42

```



```

43  if( recvfrom(sd,buf,BUFSZ,0,(struct sockaddr*)&srvr,
44      &namesize) ==-1) {
45      perror("clnt2.c:main:recvfrom");
46      exit(1);
47  }
48  printf("CLIENT REC'D: %s\n", buf);
49  close(sd);
50  unlink( clnt.sun_path );
51  }

```



## AF\_INET, SOCK\_STREAM Example

### Server Side

```

1  /* srvr3.c - Server - full example
2  * AF_INET, SOCK_STREAM Time Server
3  * Server sends its current time */
4
5  #include <stdio.h>
6  #include <sys/types.h>
7  #include <sys/socket.h>
8  #include <sys/utsname.h>
9  #include <netdb.h>
10 #include <netinet/in.h>
11 #include <errno.h>
12 #include <time.h>
13
14 #define PORTNUM 5998
15

```



```

16  main() {
17
18      struct utsname name;
19      struct sockaddr_in socketname, client;
20      int sd, ns, clientlen = sizeof(client);
21      struct hostent *host;
22      time_t today;
23
24      /* determine server system name and internet address */
25      if (uname(&name) == -1) {
26          perror("uname");
27          exit(1);
28      }
29
30      if ((host = gethostbyname(name.nodename)) == NULL) {
31          perror("gethostbyname");
32          exit(1);
33      }
34

```



```

35  /* fill in socket address structure */
36  memset((char *) &socketname, '\0', sizeof(socketname));
37  socketname.sin_family = AF_INET;
38  socketname.sin_port = PORTNUM;
39  memcpy( (char *) &socketname.sin_addr,
40          host->h_addr,host->h_length);
41
42  /* open socket */
43  if ((sd = socket(AF_INET, SOCK_STREAM, 0)) == -1) {
44      perror("socket");
45      exit(1);
46  }
47
48  /* bind socket to a name */
49  if (bind(sd, (struct sockaddr *) & socketname,
50          sizeof(socketname))) {
51      perror("bind");
52      exit(1);
53  }
54

```



```

55  /* prepare to receive multiple connect requests */
56  if (listen(sd, 128)) {
57      perror("listen");
58      exit(1);
59  }
60
61  while (1) {
62      if ((ns = accept(sd, (struct sockaddr *)&client,
63                      &clientlen)) == -1) {
64          perror("accept");
65          exit(1);
66      }
67
68      /* get current time and date */
69      time(&today);
70

```



```

71  /* send time through socket */
72  if (write(ns, &today, sizeof(today)) == -1) {
73      perror("write");
74      exit(1);
75  }
76  close(ns);
77  }
78  }

```



## Client Side

```

1  /* clnt3.c - Client full example
2  * AF_INET, SOCK_STREAM Client of Time Server */
3
4  #include <stdio.h>
5  #include <sys/types.h>
6  #include <sys/socket.h>
7  #include <sys/utsname.h>
8  #include <netdb.h>
9  #include <netinet/in.h>
10 #include <errno.h>
11 #include <time.h>
12
13 #define PORTNUM 5998
14
15 main() {
16
17     int sd;
18     struct sockaddr_in server;
19     struct hostent *host;
20     time_t srvrtime;

```



```

21
22     struct utsname name;
23     /* create the socket for talking to server*/
24     if ((sd = socket(AF_INET, SOCK_STREAM, 0)) == -1) {
25         perror("socket");
26         exit(1);
27     }
28
29     /* get server internet address and put into addr
30      * structure fill in the socket address structure
31      * and connect to server */
32     memset((char *) &server, '\0', sizeof(server));
33     server.sin_family = AF_INET;
34     server.sin_port = PORTNUM;
35
36     /* Server is local system. Get its name. */
37     if (uname(&name) == -1) {
38         perror("uname");
39         exit(1);
40     }
41     if ((host = gethostbyname("name.nodename")) == NULL) {
42         perror("gethostbyname");

```



```
43     exit(1);
44 }
45 memcpy((char *)&server.sin_addr, host->h_addr,
46         host->h_length);
47
48 /* connect to server */
49 if( connect(sd, (struct sockaddr *)&server,
50           sizeof(server)) {
51     perror("connect");
52     exit(1);
53 }
54
55 /* read the time and date passed from the server */
56 if(read(sd, &srvrtime, sizeof(srvrtime) ) == -1) {
57     perror("recv");
58     exit(1);
59 }
60 close(sd);
61
62 printf("Server's time is: %s", ctime(&srvrtime));
63 }
```



## Exercise: Sockets

- Objectives
- Tasks
- Discussion
- Solutions